

Scientific Data Analysis Pipelines -

Push, Pull, React, Or Schedule?

Ami Tavory
Final, Israel

SciPy.in 2012

PyConTW 2013


Outline



- Introduction
- Usage Consideration
- Available Language Constructs
- Push, Pull, React, Or Schedule?
- DagPype – A Push-Based Solution
- Conclusions

Complementing Existing Python Scientific Libraries

- Once data is assembled, existing Python Scientific libraries are great
 - numpy/scipy, pandas, scikit-learn, mayavi, many more



- Data-analysis processes include several complementary tasks:
 - Data assembly (ugly-format files, outliers, etc.)
 - Extremely quick preliminary analysis
 - Huge-data aggregation (from simple correlation to complex stream algs)

Complementary Tasks: Examples And Attributes

- Examples of common tasks:
 - Find the correlation of the (transformed) content of two files, excluding outliers

wind.dat

10.2

11.2

~~9.1~~

8.3

~~840.88~~

12

13.4

...

rain.dat

0.2

0.14

~~11.4~~

0

~~0.23~~

1.2

0.01

...



correlation: 0.23

Complementary Tasks: Examples And Attributes

- Examples of common tasks:
 - Find the correlation of the contents of two files, excluding outliers
 - Aggregate similar-row CSV columns

meteo.csv

day,wind,rain,snow

1,10.2,3.8,4.7

1, 12.2,4.8

1,10.7,6.7

1,9.7,5.2

2,10.4,8.77

2,9,6.2

...



[9.97,10,...]

Complementary Tasks: Examples And Attributes

- Examples of common tasks:
 - Find the correlation of the contents of two files, excluding outliers
 - Aggregate similar-row CSV columns
- Common attributes:
 - Complex & flexible combination of relatively simple steps
 - (Existing libraries focus on complex steps)

Pipes

- Pipes are natural for expressing multiple stage connections.
 - This is how we use the shell:

```
$ cat | grep | wc
```

Pipes

- Pipes are natural for expressing multiple stage connections.
- Previous two examples, using extended pipeline syntax:

```
>> wind_rain_corr = stream_vals('wind.txt') +  
stream_vals('rain.txt') | \  
... filt(pre = lambda (wind, rain) : wind < 10 and rain < 10)  
| \  
... corr()  
  
>> day_wind = stream_vals('meteo.csv', ('day', 'wind')) | \  
... group( \  
...     key = lambda (day, wind, rain, snow) : day,  
...     pipe = lambda day : ave_()) | \  
... to_array()
```

valid Python expressions can be used in scripts or even
a(n IronPython) prompt.

Plenty Of Reusable Stages

IO

stream_lines (text), stream_vals (csv),
parse_xml (xml),
to_stream (text), vals_to_stream (csv),
np.chunk_stream_vals (csv, numpy), ...

Aggregates

min_, max_, sum, count, mean,
stddev, corr, kurtosis

Econometrics

...

Control

filt, skip, nth, cast, count,
to_list, to_dict, np.to_array,
size_rand_sample, prob_rand_sample
select_inds, grep, trace

Signal Processing


window_simple_ave, cum_ave, exp_ave,
window_min, window_max, window_quantile
Plotting ...

...

Underlying Mechanism?


- Which underlying mechanism should drive the pipes?
 - David Beazly (Pycon 08, 09)
- Focus of this talk
- We will:
 - Look at usage considerations
 - Review available language constructs
 - See, using these points, why a coroutine *push*-based solution is necessary

Outline

- Introduction
-  • Usage Consideration
- Available Language Constructs
- Push, Pull, React, Or Schedule?
- DAGPytype – A Push-Based Solution
- Conclusions

Footprint

- Typical data sets can be huge.
- Each stage should be able to process the data in chunks, not necessarily its entirety.
- Memory use should be proportional to the size of the chunks, not the size of the original dataset.
 - First example (correlation): constant




Footprint

- Typical data sets can be huge.
- Each stage should be able to process the data in chunks, not necessarily its entirety.
- Memory use should be proportional to the size of the chunks, not the size of the original dataset.
 - First example (correlation): constant
 - Second example (grouping): proportional to the number of days.




Topology - DAGs (Directed Acyclic Graphs)

- Traditional pipelining is linear:
 - Each stage is fed by at most one stage, and feeds at most one stage.



- Data-processing pipelines need more expressiveness.



- Fanning in, fanning out
- In general, a DAG (directed acyclic graph) topology is needed.

```
source0() + (source1() | filt0()) | \  
  filt1() | \  
  sink0() + sink1() + (filt2() | sink2())
```


Python Interoperability

- Pipelines are natural for expressing some things, not all
 - Loops, conditionals, functions, classes
- Pipes should work with the rest of the language

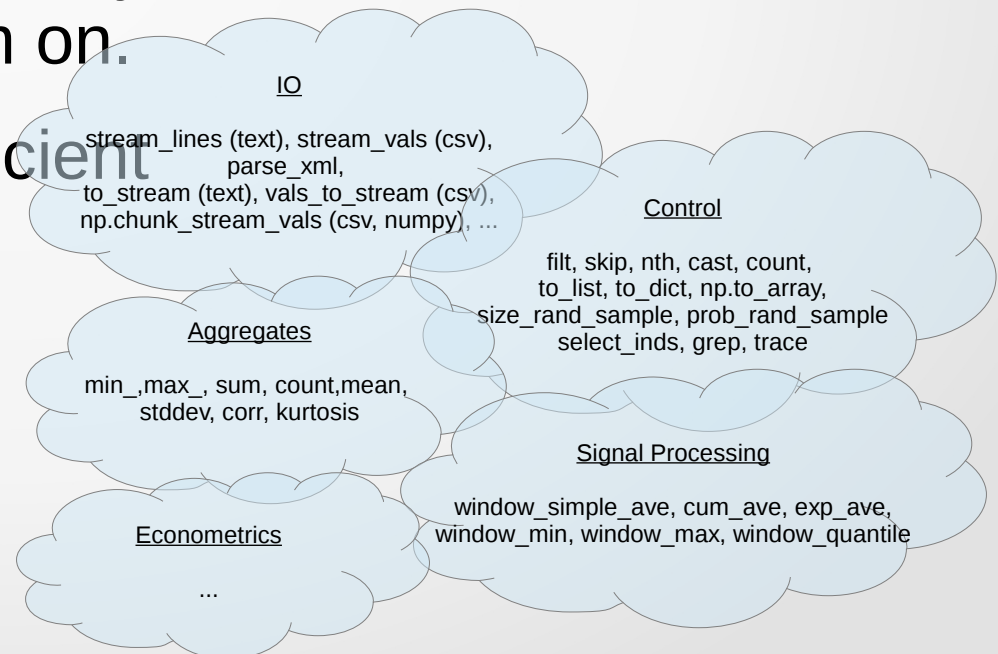
```
means = [ stream_vals( '%d_data.dat' % i ) . mean( ) for i in  
xrange( 10 ) ]
```

Relay / Processing Ratio

- Python pipelines used for many tasks, e.g., web servers



- Typically have low relay / processing ratio
- Conversely, stages here typically take a few floats, do some operations, send them on.
- Relay should be direct & efficient



Stage State And Synchronization

- In general, we must assume that processing stages are stateful...

```
wi nd_r ai n_c or r  = st ream_val s( 'm et eo. csv', ( 'wi nd', 'r ai n' )) |  
cor r ( )
```

(x, y) → **corr** {n, sum_xtx, sum_xty, sum_yty}

- and unsynchronized

```
st ream_val s( 'm et eo. csv', 'r ai n' ) | r el ay ( ) + ski p ( 5 ) | cor r ( )
```





Outline

- Introduction
- Usage Consideration
- ➔ • Available Language Constructs
- Push, Pull, React, Or Schedule?
- DagPype – A Push-Based Solution
- Conclusions

Starting Point - Statefulness

- As we saw, when stages process elements, they change state



- Which stateful mechanisms does Python offer?

Class Objects - *React*

- Classic OO:
 - State stored in objects
 - Objects *react* through methods

```
1 class Kl ass( obj ect ) :  
2     def __i ni t __ ( self , .. ) :  
3         ... [ St ore i ni ti al st at e ] ...  
4  
5     def met hod ( self , ... ) :  
6         ... [ Updat e st at e ] ...  
7  
8     def anot her _met hod ( self , ... ) :  
9         ... [ Updat e st at e ] ...
```

Generators - *Pull*

- Writing style sequential, but execution not

```
1 def gen(some_sequence, ...):  
2     ... [set initial state] ...  
3     for e in some_sequence:  
4         ... [update state ...]  
5         yield ... [something depending on s] ...
```

like calling a function

- Example: transform
cumulative sum.

like returning a value ... but not quite

– [1, 2, 3, 4] → [1, 3, 6, 10]

```
1 def cum_sum(seq):  
2     sum_ = 0  
3     for e in seq:  
4         sum_ += e  
5         yield sum
```

Coroutines - *Push*

- Again, writing style sequential, but execution not

```
1 def co(target, ...):  
2     ... [set initial state] ...  
3     try:  
4         while True:  
5             e = (yield)  
6             target.send(... [something] ...)  
7     except GeneratorExit:  
8         ... [send something]  
9         target.close()
```

like being called ... but not quite

like calling a function

```
1 def cum_sum(target):  
2     sum = 0  
3     try:  
4         while True:  
5             e = (yield)  
6             sum += e  
7             target.send(sum)  
8     except GeneratorExit:  
9         target.close()
```

Trampoline-Function Coroutines - *Schedule*

- Using coroutines yielding functions, it's possible to build a full-blown scheduler.
- Won't go into it.

Outline

- Introduction
- Usage Consideration
- Available Language Constructs
- ➡ • Push, Pull, React, Or Schedule?
- DagPy – A Push-Based Solution
- Conclusions

React-Style Stages

- A simple absolute-value stage:
 - $[1, -2, -3, 4] \rightarrow [1, 2, 3, 4]$

```
1  class Abs(object):
2      def __init__(self):
3          self._has_val = False
4
5      def push(self, e):
6          self._has_val = True
7          self._val = abs(e)
8
9      def close(self):
10         pass
11
12     def has_value(self):
13         return self._has_val
14
15     def value(self):
16         assert self._has_val
17         self._has_val = False
18         return self._val
```

needs to be set somewhere

push a value, update one exists

some stages perform after-close ops

stages can be unsynchronized

get a value (assuming one exists)

React-Style Stages

- A simple absolute-value stage:

```
1  class Abs( object ):
2      def __init__( self ):
3          self._has_val = False
4
5      def push( self, e ):
6          self._has_val = True
7          self._val = abs( e )
8
9      def close( self ):
10         pass
11
12     def has_value( self ):
13         return self._has_val
14
15     def value( self ):
16         assert self._has_val
17         self._has_val = False
18         return self._val
```

– Code SNR, Runtime

Pull-Style Stages

- A simple absolute-value stage:

```
1 @pi pe
2 def abs( seq):
3     for e in seq:
4         yield abs( e)
```

Advantages Of *Pull*-Style Stages

- Simplicity, low overhead

```
@pipe
def abs( seq ):
    for e in seq:
        yield abs( e)
```

Advantages Of *Pull-Style* Stages


- Simplicity, low overhead
- “Free” left associativity

```
stream_val s('foo.dat') | abs_() | mean()
```


```
1 @pipe
2 def abs_(seq):
3     for e in seq:
4         yield abs(e)
```

Major Disadvantages Of *Pull-Style Stages*

- No efficient fan-out
 - Requires memory loading entire stream

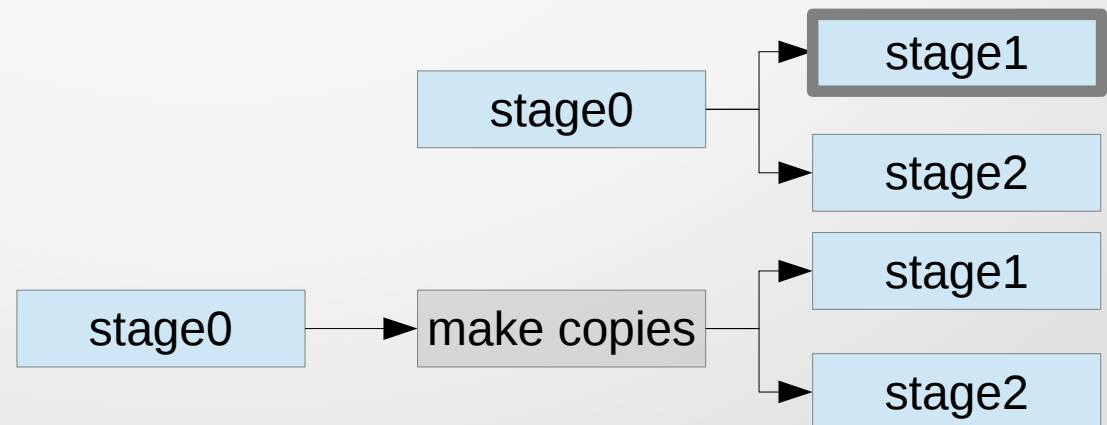


- Implication – no composable aggregates



- Inherent nature of generators

```
for e in seq:  
    ... [do something] ...
```



Push-Style Stages

- A simple absolute-value stage:

```
1 @filter
2 def abs(target):
3     try:
4         while True:
5             target.send(abs(yield))
6     except GeneratorExit:
7         target.close()
```

Assessment Of *Push*-Style Stages


- Clarity and overhead slightly less than *pull*, but much better than *react*
- Natural associativity is right, not left

```
stream_vals('foo.dat') | abs_() | mean()
```

```
1 @filter
2 def abs(target):
3     try:
4         while True:
5             target.send(abs((yield)))
6     except GeneratorExit:
7         target.close()
```


Assessment Of *Push*-Style Stages

- Clarity and overhead slightly less than *pull*, but much better than *react*
- Natural associativity is right, not left
- Slightly unnatural fan-in



```
1  def stage2( target )
2      try:
3          while True:
4              x, y = (yield)
5              ...
6          except GeneratorExit:
7              ...
8              target.close()
```

Outline

- Introduction
- Usage Consideration
- Available Language Constructs
- Push, Pull, React, Or Schedule?
- DAGPy – A Push-Based Solution
 - Some Design Points
 - Examples & Performance
- Conclusions



Lazy Construction

- LTR / RTL discrepancy solution
- During pipeline construction, objects store pipes and fans recursively as lists and tuples.

```
source0() + (source1() | filter0()) | \
    filter1() | \
    sink0() + sink1() + (filter2() | sink2())
```

```
[source0]
```

```
[source1]
```

```
[filter0]
```

Lazy Construction

- LTR / RTL discrepancy solution
- During pipeline construction, objects store pipes and fans recursively as lists and tuples.

```
source0() + (source1() | filter0()) | \
filter1() | \
sink0() + sink1() + (filter2() | sink2())
```

```
[source0]
```

```
[source1, filter0]
```

Lazy Construction

- LTR / RTL discrepancy solution
- During pipeline construction, objects store pipes and fans recursively as lists and tuples.

```
source0() + (source1() | filter0()) | \
    filter1() | \
    sink0() + sink1() + (filter2() | sink2())
```

```
[(source0, [source1, filter0])]
```

Lazy Construction

- LTR / RTL discrepancy solution
- During pipeline construction, object store pipes and fans recursively as lists and tuples.


```
source0() + (source1() | filt0()) \
  filt1() | \
  sink0() + sink1() + (filt2() | sink2())
```

```
[(source0, [source1, filt0]), filt1, (sink0, sink1, [filt2, sink2])]
```

- Source->sink connections trigger actual calculations.

Fanning Out And In

- Fanning out:





```
stage0 | stage1 + stage2
```

```
1 def stage0(target):  
2     while True:  
3         e = (yield)  
4         target.send( fn(e) )
```

Fanning Out And In


- Fanning out:




```
1 def _demux(*targets):
2     ...
3     while True:
4         e = (yield)
5         for t in *targets:
6             t.send(e)
```


Fanning Out And In

- Fanning out:




- Fanning in:




```
1 def stage2(target):  
2     ...  
3     while True:  
4         x, y = (yield)
```

Fanning Out And In

- Fanning out:




- Fanning in:



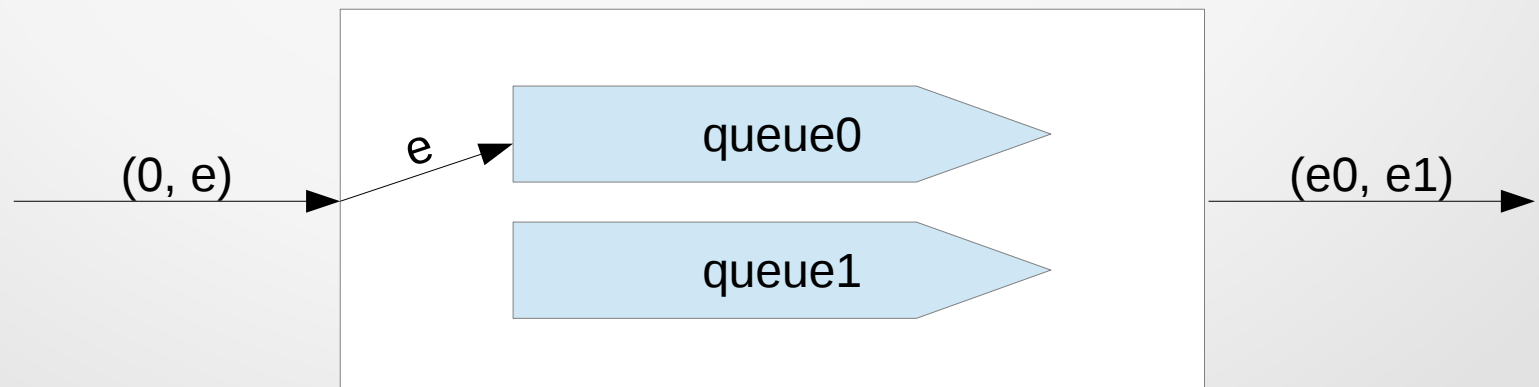

```
1 def _id(target, id):  
2     ...  
3     while True:  
4         target.send( (id, (yield)) )
```

Fanning Out And In

- Fanning out:



- Fanning in:



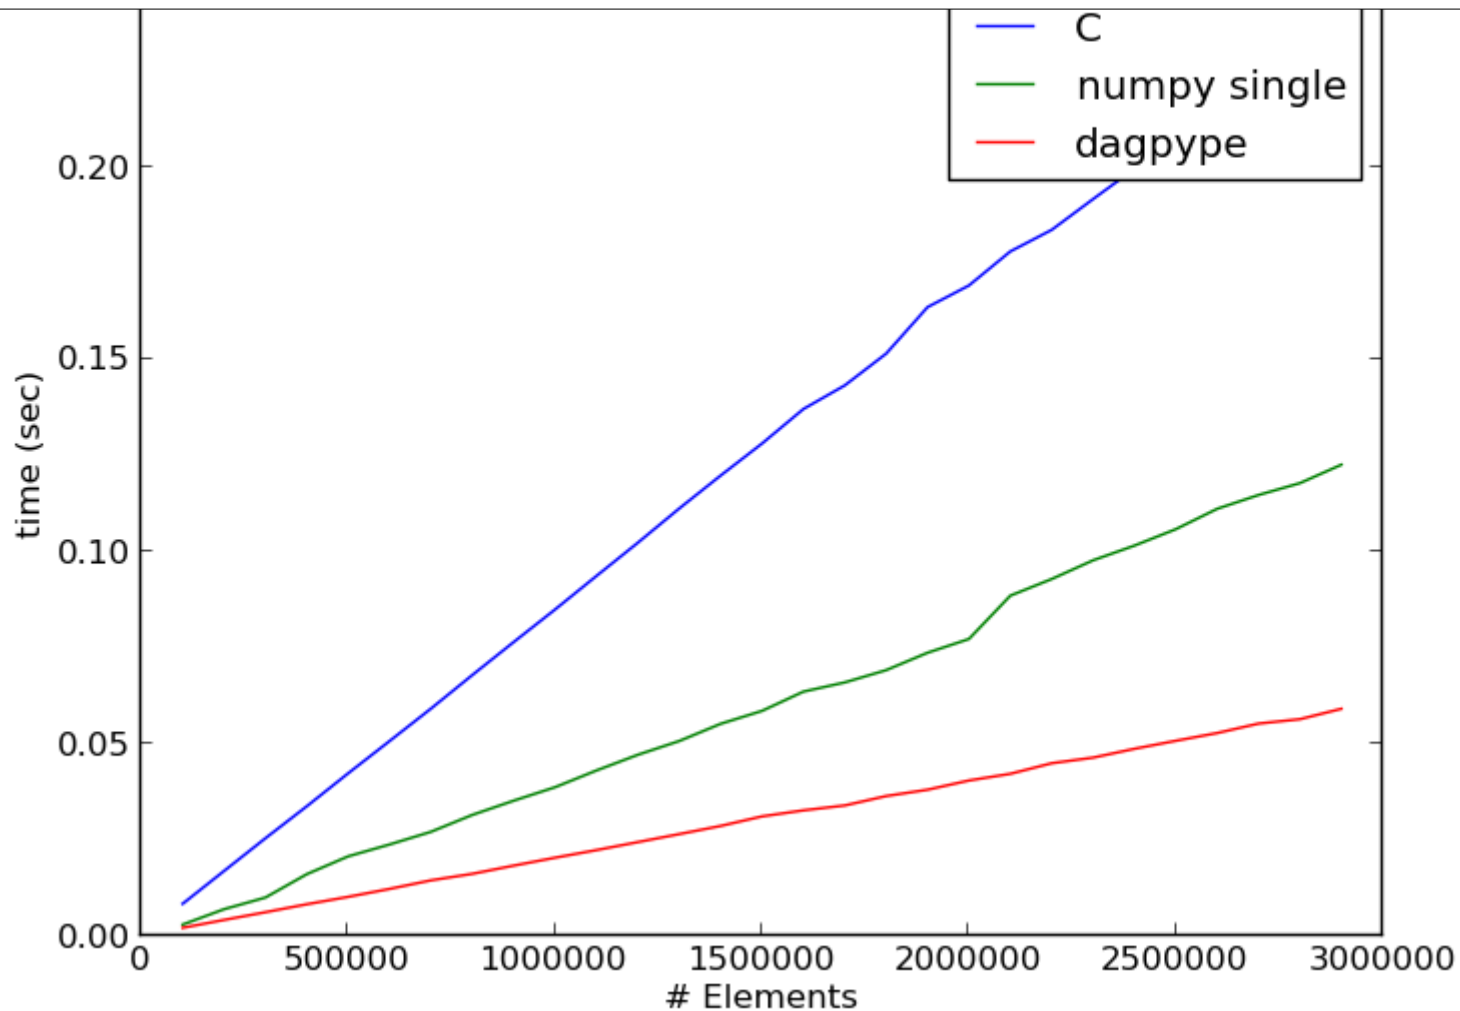
Outline

- Introduction
- Usage Consideration
- Available Language Constructs
- Push, Pull, React, Or Schedule?
- DagPype – A Push-Based Solution
 - Some Design Points
 - Examples & Performance
- Conclusions




Performance - Binary File Correlation

```
c = np.chunk_stream_bytes('foo.dat', num_cols = 2) | np.corr()
```

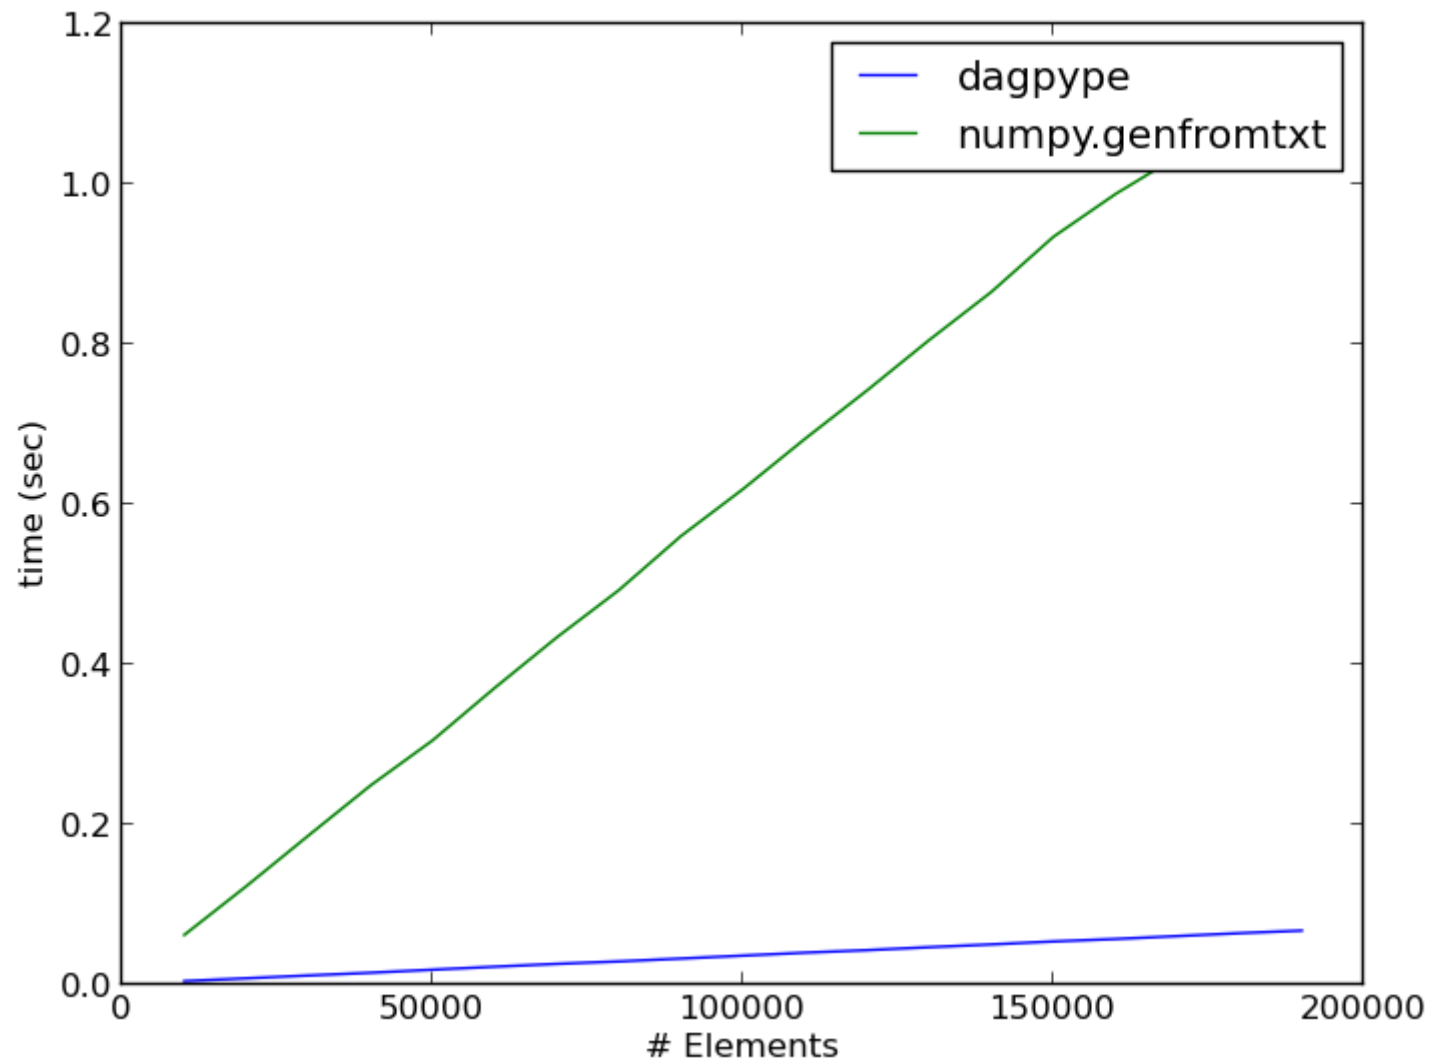


Performance - Pruned Binary File Correlation

```
c = np.chunk_stream_bytes('foo.csv', num_cols = 2) | \
    filt(lambda a : a[numpy.logical_and(a[:, 0] < 0.25, a[:, 1] <
0.25), :]) | \
    np.corr()
```



Performance - CSV File Mean



Outline

- Introduction
- Usage Consideration
- Available Language Constructs
- Push, Pull, React, Or Schedule?
- DagPype – A Push-Based Solution
- ➔ • Conclusions

Main Points Covered

- (Python) pipelines can complement existing great Python libraries and tools
 - A large fraction of (at least my) time is spent
 - trying to figure out what to learn ...
 - out of horrible (noisy) data ...
 - in terrible formats

Main Points Covered

- (Python) pipelines can complement existing great Python libraries and tools
- *Push*-based coroutines are the way to implement them in Python

<https://pypi.python.org/pypi/DAGPype/>

Points Not Covered

- Chunk granularity
- Parallelism
- Operations beyond DAG-pipelines' expressiveness

Thanks!

- Thank you for your time!

